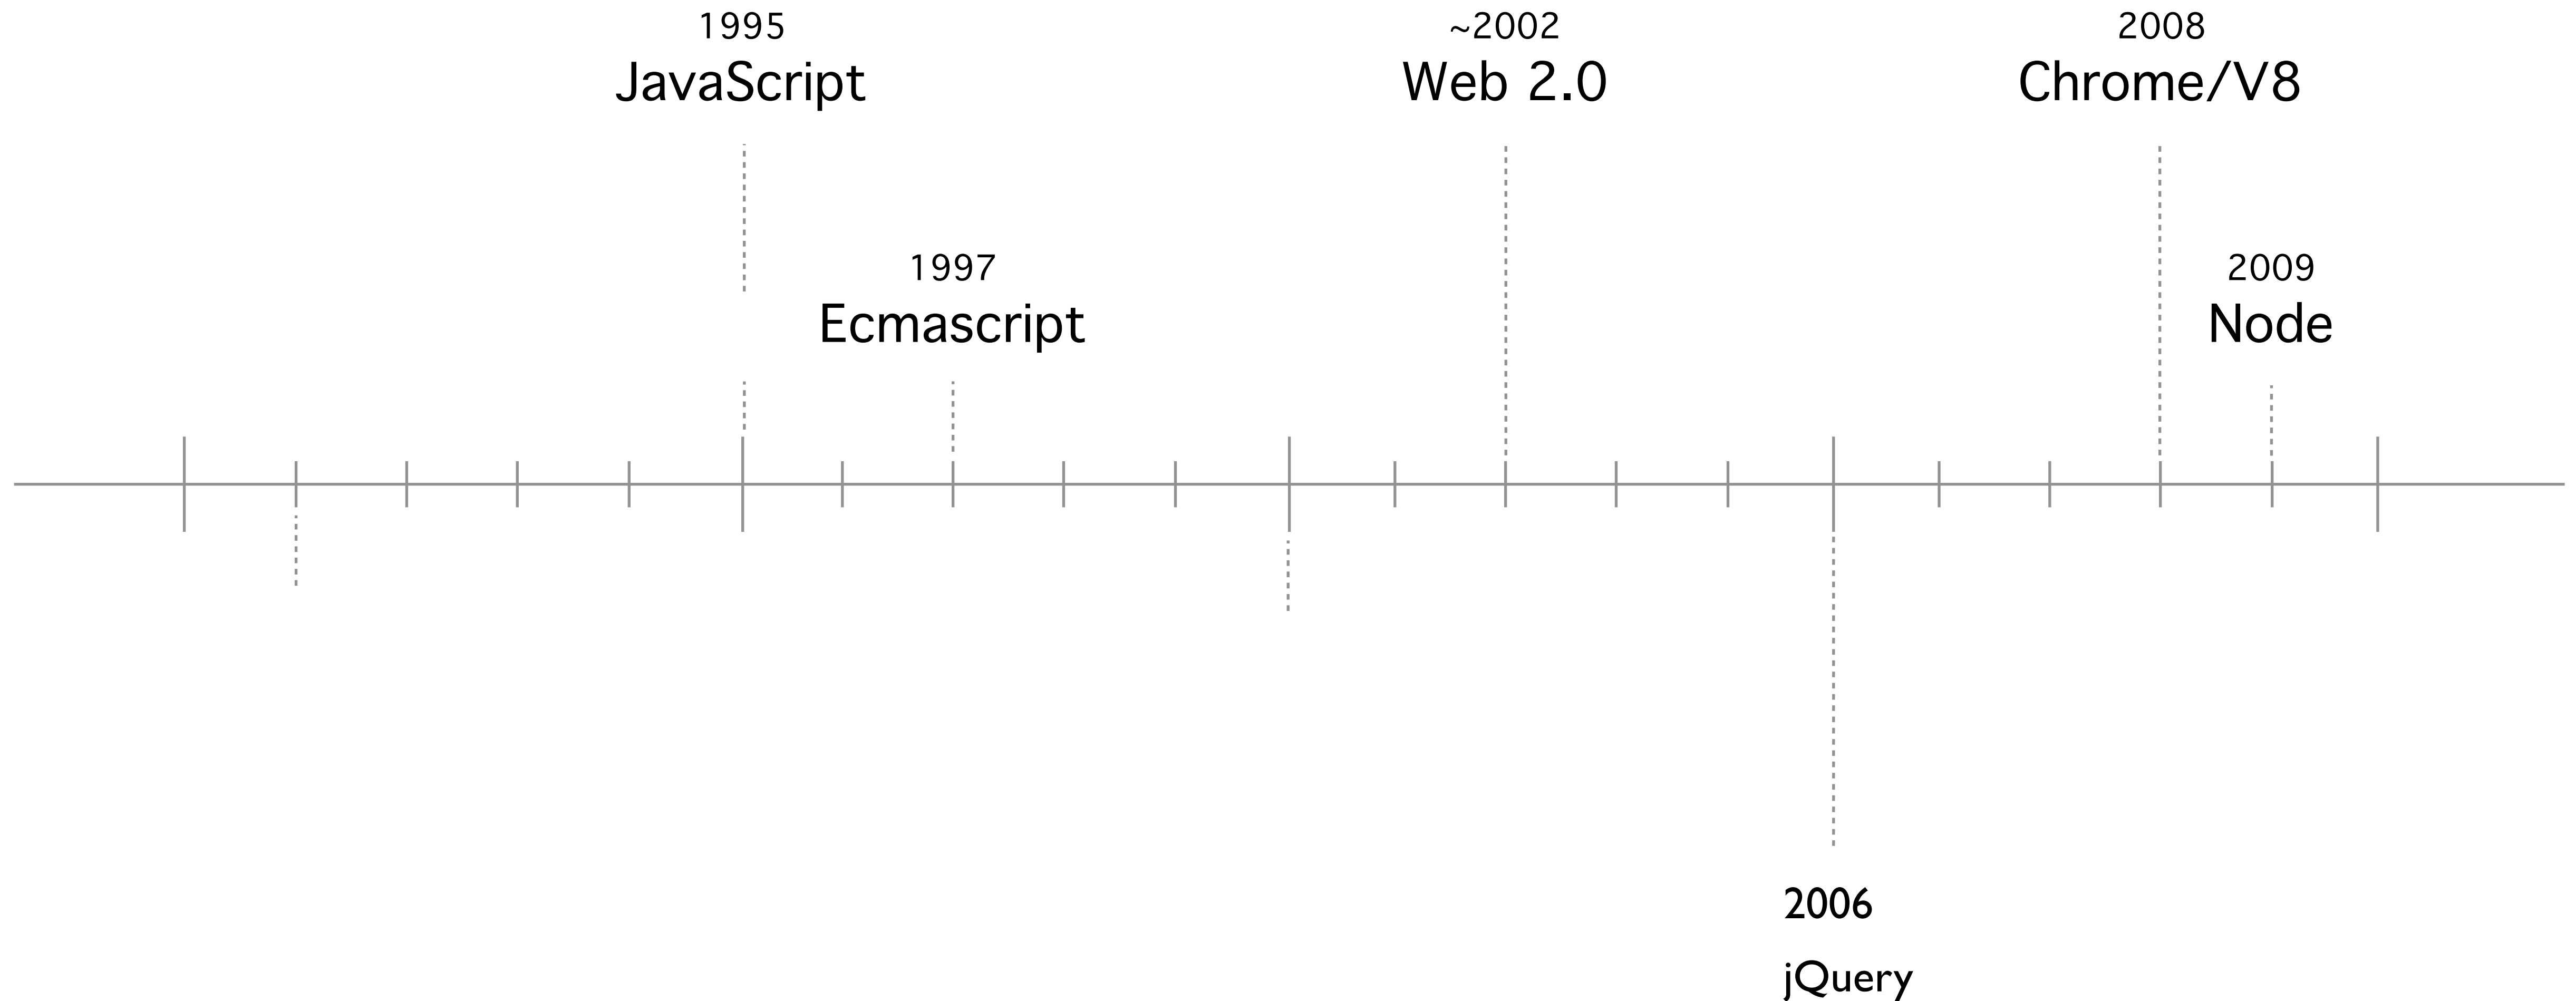






A Brief History of Everything



Intro to



or: Kernals, Processes, Threads — Oh My!
or: Fun with Modules
or: The Kitchen Sink & Async
or: Six Degrees from Ryan Dahl

Outline

- **What is Node?**
- **Program vs Process**
- **History / context**
- **Node as process**
- **Node modules**
- **Node asynchronicity**

What is Node?

“Node.js® is a JavaScript runtime built on **Chrome's V8 JavaScript engine**. Node.js uses an event-driven, non-blocking I/O model that makes it lightweight and efficient. Node.js' package ecosystem, **npm**, is the largest ecosystem of open source libraries in the world.”

Node.js: the Essentials

- Node is a program (written in C, C++, and JS)
- It can run JavaScript via the Chrome V8 engine (C++)
- It provides APIs for your OS's file / network system (slow I/O)
- In other words, Node lets you run JS on a computer *outside of a browser* and interact with stuff on that computer
- Why?



The Practical Dev

@ThePracticalDev

Follow



So NASA uses Node.js in their space-suits.
The logical follow up question: Do NASA
space-suits depend on left-pad?



Benjamin Coe @BenjaminCoe

6/3/16

@CollinEstes I heard a rumor that NASA
uses Node.js for space-suits. I'm
curious, do you use the npm ecosystem
to develop these apps?



57

69



Collin Estes @CollinEstes


6/3/16

@BenjaminCoe You heard correctly, and
yes we do.



36

74



Program vs. Process

- **Program is data**
 - machine code (pre-compiled)
 - bytecode
 - text file (can be interpreted)
- **Inert — not doing anything**
- **Ready to be run as a process**
- **Process is execution**
 - memory allocated
 - CPU performing steps
- **"Live"**
- **Produces results**
- **Interactive**
- **Can be started/stopped**
- **Multiple processes from one program...**

Node as Process

(Demo)



Different Environments, Same Engine



fs

process

net



window

history

document

Browser Tool Examples

- **Synchronous stuff:**
 - console
 - window
 - document
- **Asynchronous stuff:**
 - `setTimeout(callback, delay)` (note: not in ECMAScript, "not JS" — an API!)
 - `setInterval(callback, delay)` — ditto above
 - `new XMLHttpRequest()` / `.open`, etc. — making HTTP requests
 - `element.onclick(callback)` // slightly different example

Node Tool Examples

- **Synchronous stuff:**

- console (same name, but a wrapper for process.stdout)
- global (no window!)
- module (we'll get to this)

- **Asynchronous stuff:**

- setTimeout(callback, delay) — same name for convenience's sake
- setInterval(callback, delay) — ditto
- http.request(options, callback) — uses built-in `http` JS module
- etc.

Modularity

What is the purpose of a modular system?

- Organized and maintainable vehicle for code splitting
- Why separate code?
 - Single files have more defined responsibility
 - Easier collaboration
 - Visible structure
 - Testing
 - Maintainability
 - Reusability

Built-in Modules

- There are MANY, here are a few:
 - url
 - path
 - fs
 - http
 - crypto
 - net
 - child_process

How do we get more?

- Author-defined
- Third-party



Asynchronous / Non-blocking I/O

Big Use Case: Server


- A program running on a computer
- Listens on a *port* for *requests*
- Sends *responses* back
- Communication must follow agreed-upon *protocol* (e.g. HTTP)
- [DEMO]
- Problem: communication over network is super-slow! Should be able to field other requests / do other stuff in meantime.

Slow *blocking* I/O is handed off to a thread


- CPU clock cycles per access level (*Dahl 2010):
 - Dynamic memory: "non-blocking"
 - L1 cache: 3
 - L2 cache: 14
 - RAM: 250
 - "Blocking" I/O:
 - Disk: 41,000,000
 - Network: 240,000,000

Concurrency


blocking



non-blocking



parallel (blocking)



"Threads are Hard"

- Every thread comes with performance overhead — memory, CPU, etc.
- Switching threads within a process consumes clock cycles (have to schedule jumps)
- Switching threads between processes consumes clock cycles AND waits for memory (have to load process state)
- Communicating between threads in the same process is very tricky — synchronizing results, being efficient
- Thread programming best left to experts

Server Strategies



- Apache: every new connection gets its own thread (wasteful of resources, doesn't scale well, lots of work to synchronize)



- Nginx: an event loop which queues *callbacks* back into the main execution stack, once thread managing blocking code finishes

*“Browsers got it right — abstract that to
callbacks with the DOM API.”*

RYAN DAHL, CREATOR OF NODE

Er, not exactly

*“Node.js is a ~~single threaded~~, event-driven,
non-blocking I/O platform”*

– SOME PEOPLE ON THE INTERNET

“JavaScript is single-threaded” ...arguably yes

– OTHER PEOPLE ON THE INTERNET



JS

Standard Library (modules)
(e.g. fs, http, etc.)

C &
C++

C bindings (glue)

V8




Thread
pool
(libeio)

Event
loop
(libev)

Also: secure crypto,
DNS... but this is the
core functionality

Event Loop

JavaScript:
One thread



Thread pool (libeio):
Slow stuff, multiple
threads

Event loop (libev):
One thread